

1 Final Project Submission: Real Estate Linear Regression Model Analysis

(Phase 2)

- Student Name: Tenicka Norwood
- Program Pace: self paced
- Scheduled Project Review Time:
- Instructor name: Joe Comeaux
- Blog post Url: <https://medium.com/@tenicka.norwood/working-with-data-while-trying-to-stay-dry-fa4ebf5e5f64>
(<https://medium.com/@tenicka.norwood/working-with-data-while-trying-to-stay-dry-fa4ebf5e5f64>)



Photo by: anyaberkut on [Canva](https://www.canva.com) (<https://www.canva.com>)

2 Overview:

LandingPad Realtors is a real estate business that helps families with school-aged children relocate to King County and find the perfect home to meet their families needs. LandingPad provides potential homeowners with home purchase options within their ideal budget.

For this project, I will start by identifying the characteristics of homes that increase housing costs. The effect of each relevant feature will then be identified and communicated to the team at LandingPad. This project will be grounded in performing a statistical analysis of the price of houses in the King County House dataset and creating a multiple linear regression model that accurately predicts the sale price of a house in King County.

Linear Regression is often described as a predictive model based on the sum of weighted independent variables. This mathematical linear relationship between an dependent variable and one or more dependent variables and is often shown as:

$$y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_n x_n$$

where β_0 is a constant that shows the y-intercept and β_1 to β_n are coefficients that show how the dependent variable y changes with the independent variables x_1 to x_n .

3 Business Understanding

- **Stakeholder:** LandingPad Realtors
- **Business Case:** I have been hired by LandingPad to accurately predict the housing prices within the King County Housing Market. Executives at LandingPad want to launch a multimedia campaign to reach their target audience of young families moving to the Kings County Area and want a reliable model that can be refined over time as more information becomes available.

4 Objectives

We will use the CRISP DM model to:

- Understand the Data
 - Examine and document surface properties of the data
 - Dig deeper into the data to visualize and identify relationships among the data
- Prepare the Data

- Select, clean, construct, integrate and format data
- Model the Data
 - Determine which algorithm to try
- Evaluate the Model
 - Determine if the model meets the business success criteria and determine next steps

Then, I will use this to build and refine a **linear regression model** that I can use to answer the following guiding questions:

- 1. Which neighborhoods have the highest average home price?**
Understanding the effect of neighborhood location on home price is key information for potential home owners and realtors
- 2. How does the number of bedrooms affect the sale price of a home?**
Insights on the affect of attributes on the sale price can help new home owners budget appropriately
- 3. How does proximity to a highly rated school affect the sale price of a home?**
Knowing which homes are connected to highly rated schools is vital information for families with school aged children.

The recommendations garnered by answering these questions will be valuable to LandingPad Realtors because they will help prospective home buyers confidently determine which homes yield the best options within their price range.

5 Data Understanding

In this project I will use the CRISP DM method. The dataset selected in this project are from the :

- King County House Sales Dataset found in `kc_house_data.csv`

The dataset can be found in the data folder of this repository along with a file called `column_names.md` which provides description of the features within the dataset. More information about the features on the site of [the King County Assessor. \(https://info.kingcounty.gov/assessor/esales/Glossary.aspx?type=r\)](https://info.kingcounty.gov/assessor/esales/Glossary.aspx?type=r)

The King County House Sales Dataset includes sales data for 21,597 homes with 20 features:

Name	Description	Final Datatype	Numeric or Categorical	Target or Feature
id	Unique identifier for a house	int	Numeric	Feature
date	Date house was sold	datetime	Numeric	Feature
price	Sale price (prediction target)	int	Numeric	Target
bedrooms	Number of bedrooms	int	Numeric	Feature
bathrooms	Number of bathrooms	float	Numeric	Feature
sqft_living	Square footage of living space in the home	int	Numeric	Feature
sqft_lot	Square footage of the lot	int	Numeric	Feature
floors	Number of floors(levels) in house	float	Numeric	Feature
waterfront	Whether the house is on a waterfront	float	Categorical	Feature
view	Quality of view from house	float	Categorical	Feature
condition	How good the overall condition of the house is. Related to the maintenance of house	int	Numeric	Feature
grade	Overall grade of the house. Related to the construction and design of the house	int	Categorical	Feature
sqft_above	Square footage of house apart from basement	int	Numeric	Feature
sqft_basement	Square footage of the basement	float	Numeric	Feature
yr_built	Year when house was built	int	Numeric	Feature
yr_renovated	Year when house was renovated	int	Numeric	Feature
zipcode	ZIP Code used by the United States Postal Service	int	Categorical	Feature
lat	Latitude coordinate	float	Numeric	Feature
long	Longitude coordinate	float	Numeric	Feature
sqft_living15	The square footage of interior housing living space for the nearest 15 neighbors	int	Numeric	Feature
sqft_lot15	The square footage of the land lots of the nearest 15 neighbors	int	Numeric	Feature

5.1 Data Preparation

5.2 Import libraries and Visualization Packages

Importing libraries at the beginning allows access to modules and other tools throughout this project that help to make the tasks within this project manageable to implement. The main libraries that will be used within this project include:

- `pandas` : a data analysis and manipulation library which allows for flexible reading, writing, and reshaping of data
- `numpy` : a key library that brings the computational power of languages like C to Python
- `matplotlib` : a comprehensive visualization library
- `seaborn` : a data visualization library based on `matplotlib`

```
In [1]: # Import Libraries and visualization packages
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import matplotlib.ticker as ticker
import statsmodels.api as sm
import warnings
warnings.simplefilter(action='ignore', category=FutureWarning)

# Allow plots to display and be stored inline within a notebook
%matplotlib inline

# Used for working with the z-score
from scipy import stats

# Set display option to readable format
pd.set_option('display.float_format', lambda x: '%.2f' % x)
```

5.3 Load Data Using Pandas

Read in data from `kc_house_data.csv` using `.read_csv()` from the pandas library.

```
In [2]: # Read Data from Kings County House Data
df = pd.read_csv('data/kc_house_data.csv', index_col = 0)
```

Let's look at the first five rows of the Kings County Housing Data.

```
In [3]: df.head()
```

Out[3]:

	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	condition	grade	sqft_above	sqft_basement	y
id														
7129300520	10/13/2014	221900.00	3	1.00	1180	5650	1.00	NaN	NONE	Average	7 Average	1180	0.0	
6414100192	12/9/2014	538000.00	3	2.25	2570	7242	2.00	NO	NONE	Average	7 Average	2170	400.0	
5631500400	2/25/2015	180000.00	2	1.00	770	10000	1.00	NO	NONE	Average	6 Low Average	770	0.0	
2487200875	12/9/2014	604000.00	4	3.00	1960	5000	1.00	NO	NONE	Very Good	7 Average	1050	910.0	
1954400510	2/18/2015	510000.00	3	2.00	1680	8080	1.00	NO	NONE	Average	8 Good	1680	0.0	

5.3.1 Clean the Data

In order to clean the data, I typically address missing data, place holders and datatypes. This is the most important step of this project because if data is not appropriate for the model, the results will be inherently inaccurate and my model will result in lackluster predictions.

To dig deeper into the data, I will:

- Review the datatypes found within the entire dataframe
- Address duplicates, missing and placeholder data
- Address incorrect or incongruous datatypes for the model

5.3.2 Review Datatypes within Data Frame

```
In [4]: df.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 21597 entries, 7129300520 to 1523300157
Data columns (total 20 columns):
#   Column                Non-Null Count  Dtype
---  -
0   date                  21597 non-null  object
1   price                 21597 non-null  float64
2   bedrooms              21597 non-null  int64
3   bathrooms             21597 non-null  float64
4   sqft_living           21597 non-null  int64
5   sqft_lot              21597 non-null  int64
6   floors                21597 non-null  float64
7   waterfront            19221 non-null  object
8   view                  21534 non-null  object
9   condition             21597 non-null  object
10  grade                 21597 non-null  object
11  sqft_above            21597 non-null  int64
12  sqft_basement         21597 non-null  object
13  yr_built              21597 non-null  int64
14  yr_renovated          17755 non-null  float64
15  zipcode               21597 non-null  int64
16  lat                   21597 non-null  float64
17  long                  21597 non-null  float64
18  sqft_living15         21597 non-null  int64
19  sqft_lot15            21597 non-null  int64
dtypes: float64(6), int64(8), object(6)
memory usage: 3.5+ MB
```

5.3.3 Address duplicates, missing and placeholder data

```
In [5]: # Check for placeholders throughout the entire dataframe
df.isin(['?', '#', 'NaN', 'null', 'N/A', '-']).any()
```

```
Out[5]: date                False
price                    False
bedrooms                 False
bathrooms                False
sqft_living              False
sqft_lot                 False
floors                   False
waterfront               False
view                     False
condition                False
grade                    False
sqft_above               False
sqft_basement            True
yr_built                 False
yr_renovated             False
zipcode                  False
lat                      False
long                     False
sqft_living15            False
sqft_lot15               False
dtype: bool
```

```
In [6]: # Convert sqft_basement to float
# Replace placeholder ? to NAN values
df['sqft_basement'] = pd.to_numeric(df['sqft_basement'], errors="coerce")
```

```
In [7]: ► # Check for placeholders throughout the entire dataframe
df.isin(['?', '#', 'NaN', 'null', 'N/A', '-']).any()
```

```
Out[7]: date           False
price           False
bedrooms        False
bathrooms       False
sqft_living     False
sqft_lot        False
floors          False
waterfront      False
view            False
condition       False
grade           False
sqft_above      False
sqft_basement   False
yr_built        False
yr_renovated    False
zipcode         False
lat             False
long            False
sqft_living15   False
sqft_lot15      False
dtype: bool
```

Let's look at the columns within the `pandas` dataframe.

```
In [8]: ► df.columns
```

```
Out[8]: Index(['date', 'price', 'bedrooms', 'bathrooms', 'sqft_living', 'sqft_lot',
              'floors', 'waterfront', 'view', 'condition', 'grade', 'sqft_above',
              'sqft_basement', 'yr_built', 'yr_renovated', 'zipcode', 'lat', 'long',
              'sqft_living15', 'sqft_lot15'],
              dtype='object')
```

This dataset has 20 columns with 21597 rows of data. The `waterfront`, `view`, `condition`, `grade`, `sqft_basement` and `date` are object datatypes.

Let's check which columns contain null values.

```
In [9]: ► df.isnull().sum()
```

```
Out[9]: date           0
price           0
bedrooms        0
bathrooms       0
sqft_living     0
sqft_lot        0
floors          0
waterfront     2376
view           63
condition       0
grade           0
sqft_above      0
sqft_basement   454
yr_built        0
yr_renovated    3842
zipcode         0
lat             0
long            0
sqft_living15   0
sqft_lot15      0
dtype: int64
```

5.3.4 Address incorrect or incongruous datatypes for the model

Let's convert the date from a string to a datetime object.

```
In [10]: ► # Convert data column in the dataframe to datetime
df['date'] = pd.to_datetime(df['date'], format='%m/%d/%Y')
```

In [11]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 21597 entries, 7129300520 to 1523300157
Data columns (total 20 columns):
#   Column                Non-Null Count  Dtype
---  -
0   date                  21597 non-null  datetime64[ns]
1   price                 21597 non-null  float64
2   bedrooms              21597 non-null  int64
3   bathrooms             21597 non-null  float64
4   sqft_living           21597 non-null  int64
5   sqft_lot              21597 non-null  int64
6   floors                21597 non-null  float64
7   waterfront            19221 non-null  object
8   view                  21534 non-null  object
9   condition             21597 non-null  object
10  grade                 21597 non-null  object
11  sqft_above            21597 non-null  int64
12  sqft_basement         21143 non-null  float64
13  yr_built              21597 non-null  int64
14  yr_renovated          17755 non-null  float64
15  zipcode               21597 non-null  int64
16  lat                   21597 non-null  float64
17  long                  21597 non-null  float64
18  sqft_living15         21597 non-null  int64
19  sqft_lot15            21597 non-null  int64
dtypes: datetime64[ns](1), float64(7), int64(8), object(4)
memory usage: 3.5+ MB
```

In [12]: df.head()

Out[12]:

	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	condition	grade	sqft_above	sqft_basement	yr_bui
id														
7129300520	2014-10-13	221900.00	3	1.00	1180	5650	1.00	NaN	NONE	Average	7 Average	1180	0.00	195
6414100192	2014-12-09	538000.00	3	2.25	2570	7242	2.00	NO	NONE	Average	7 Average	2170	400.00	195
5631500400	2015-02-25	180000.00	2	1.00	770	10000	1.00	NO	NONE	Average	6 Low Average	770	0.00	193
2487200875	2014-12-09	604000.00	4	3.00	1960	5000	1.00	NO	NONE	Very Good	7 Average	1050	910.00	196
1954400510	2015-02-18	510000.00	3	2.00	1680	8080	1.00	NO	NONE	Average	8 Good	1680	0.00	198

From our check, the waterfront column has 2376 null values. The view column has 63 null values and the yr_renovated column has 3842 null values.

In [13]: df.describe()

Out[13]:

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	sqft_above	sqft_basement	yr_built	yr_renovated	zipcode	lat	lo
count	21597.00	21597.00	21597.00	21597.00	21597.00	21597.00	21597.00	21143.00	21597.00	17755.00	21597.00	21597.00	21597.00
mean	540296.57	3.37	2.12	2080.32	15099.41	1.49	1788.60	291.85	1971.00	83.64	98077.95	47.56	-122.00
std	367368.14	0.93	0.77	918.11	41412.64	0.54	827.76	442.50	29.38	399.95	53.51	0.14	0.00
min	78000.00	1.00	0.50	370.00	520.00	1.00	370.00	0.00	1900.00	0.00	98001.00	47.16	-122.00
25%	322000.00	3.00	1.75	1430.00	5040.00	1.00	1190.00	0.00	1951.00	0.00	98033.00	47.47	-122.00
50%	450000.00	3.00	2.25	1910.00	7618.00	1.50	1560.00	0.00	1975.00	0.00	98065.00	47.57	-122.00
75%	645000.00	4.00	2.50	2550.00	10685.00	2.00	2210.00	560.00	1997.00	0.00	98118.00	47.68	-122.00
max	7700000.00	33.00	8.00	13540.00	1651359.00	3.50	9410.00	4820.00	2015.00	2015.00	98199.00	47.78	-121.00

In [14]: df['view'].value_counts()

Out[14]:

NONE	19422
AVERAGE	957
GOOD	508
FAIR	330
EXCELLENT	317

Name: view, dtype: int64

```
In [15]: > #for col in categoricals:
#         print(df[col].value_counts(), "\n")
```

```
In [16]: > df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 21597 entries, 7129300520 to 1523300157
Data columns (total 20 columns):
#   Column                Non-Null Count  Dtype
---  -
0   date                  21597 non-null  datetime64[ns]
1   price                 21597 non-null  float64
2   bedrooms              21597 non-null  int64
3   bathrooms             21597 non-null  float64
4   sqft_living           21597 non-null  int64
5   sqft_lot              21597 non-null  int64
6   floors                21597 non-null  float64
7   waterfront            19221 non-null  object
8   view                  21534 non-null  object
9   condition             21597 non-null  object
10  grade                 21597 non-null  object
11  sqft_above            21597 non-null  int64
12  sqft_basement         21143 non-null  float64
13  yr_built              21597 non-null  int64
14  yr_renovated          17755 non-null  float64
15  zipcode               21597 non-null  int64
16  lat                   21597 non-null  float64
17  long                  21597 non-null  float64
18  sqft_living15         21597 non-null  int64
19  sqft_lot15            21597 non-null  int64
dtypes: datetime64[ns](1), float64(7), int64(8), object(4)
memory usage: 3.5+ MB
```

Next I will impute missing values in the dataframe. I replaced waterfront , sqft_basement and view null values with 0 and replace yr_renovated null values with the yr_built .

```
In [17]: > df["waterfront"].value_counts()
```

```
Out[17]: NO      19075
YES        146
Name: waterfront, dtype: int64
```

```
In [18]: > df["view"].value_counts()
```

```
Out[18]: NONE      19422
AVERAGE    957
GOOD        508
FAIR        330
EXCELLENT   317
Name: view, dtype: int64
```

```
In [19]: > df["zipcode"].value_counts()
```

```
Out[19]: 98103      602
98038      589
98115      583
98052      574
98117      553
...
98102      104
98010      100
98024       80
98148       57
98039       50
Name: zipcode, Length: 70, dtype: int64
```

```
In [20]: > df["waterfront"].fillna(0, inplace = True)
df["sqft_basement"].fillna(0, inplace = True)
df["view"].fillna(0, inplace = True)
df["yr_renovated"].fillna(df['yr_built'], inplace=True)
```

```
In [21]: ► for col in df.columns:
           print('{} : {}'.format(col, df[col].nunique()))
```

```
date : 372
price : 3622
bedrooms : 12
bathrooms : 29
sqft_living : 1034
sqft_lot : 9776
floors : 6
waterfront : 3
view : 6
condition : 5
grade : 11
sqft_above : 942
sqft_basement : 303
yr_built : 116
yr_renovated : 117
zipcode : 70
lat : 5033
long : 751
sqft_living15 : 777
sqft_lot15 : 8682
```

Now let's check to see if there are any null values in any of the columns of our dataframe.


```
In [22]: ► df.isnull().sum()
```

```
Out[22]: date      0
         price     0
         bedrooms  0
         bathrooms 0
         sqft_living 0
         sqft_lot   0
         floors     0
         waterfront 0
         view       0
         condition  0
         grade      0
         sqft_above 0
         sqft_basement 0
         yr_built   0
         yr_renovated 0
         zipcode    0
         lat        0
         long       0
         sqft_living15 0
         sqft_lot15 0
         dtype: int64
```

```
In [23]: ► df['yr_renovated'] = df['yr_renovated'].astype('int64')
```

Great, now let's check out the correlation of features within my dataframe.


```
In [24]: # Review correlations on filtered dataset
plt.figure(figsize = (15,10)).suptitle("Correlation Between Home Features")
sns.heatmap(df.corr(), annot = True, cmap = "Blues");
```



The correlation shown is called a Pearson correlation and is given by the ratio below:

$$r = \frac{\sum(x_i - \bar{x})(y_i - \bar{y})}{\sqrt{\sum(x_i - \bar{x})^2} \sqrt{\sum(y_i - \bar{y})^2}}$$

A commonly held rule for interpreting the Pearson correlation is that:

- Very Strong Correlation → 0.8 to 1
- Strong Correlation → 0.6 to 0.799
- Moderate Correlation → 0.4 to 0.599
- Weak Correlation → 0.2 to 0.399
- Very Weak Correlation → 0 to 0.199

From the heatmap, it appears that house price (the target) has the **strongest correlation** with sqft_living (0.7), a weak correlation with floors (0.26), a **strong correlation** with sqft_above (0.61), a weak correlation with lat (0.31) and a moderate correlation with sqft_living_15 (0.59). The sqft_living, bathrooms, sqft_above and sqft_living_15 features all seem to be highly correlated (> 0.7) to one another as well. This will be helpful when building our baseline model and refining it after reviewing our metrics.

Looking at a heatmap this complex may be a stretch. Let's create a more elegant way to determine which pairs of features have strong correlation values.

```
In [25]: ► _df = df.corr().abs().stack().reset_index().sort_values(0, ascending = False)
         _df['pairs'] = list(zip(_df.level_0, _df.level_1))
         _df.set_index(['pairs'], inplace = True)
         _df.drop(columns = ['level_1' , 'level_0'], inplace = True)
         _df.columns = ['cc']
         _df.drop_duplicates(inplace = True)
         _df[(_df.cc > 0.6) & (_df.cc < 1)]
```

Out[25]:

	cc
(sqft_above, sqft_living)	0.88
(sqft_living, sqft_living15)	0.76
(bathrooms, sqft_living)	0.76
(sqft_above, sqft_living15)	0.73
(sqft_lot, sqft_lot15)	0.72
(price, sqft_living)	0.70
(bathrooms, sqft_above)	0.69
(sqft_above, price)	0.61

5.3.5 Scatter matrix

Create a scatter matrix for the King County House data. (This takes awhile to run and generate a plot)

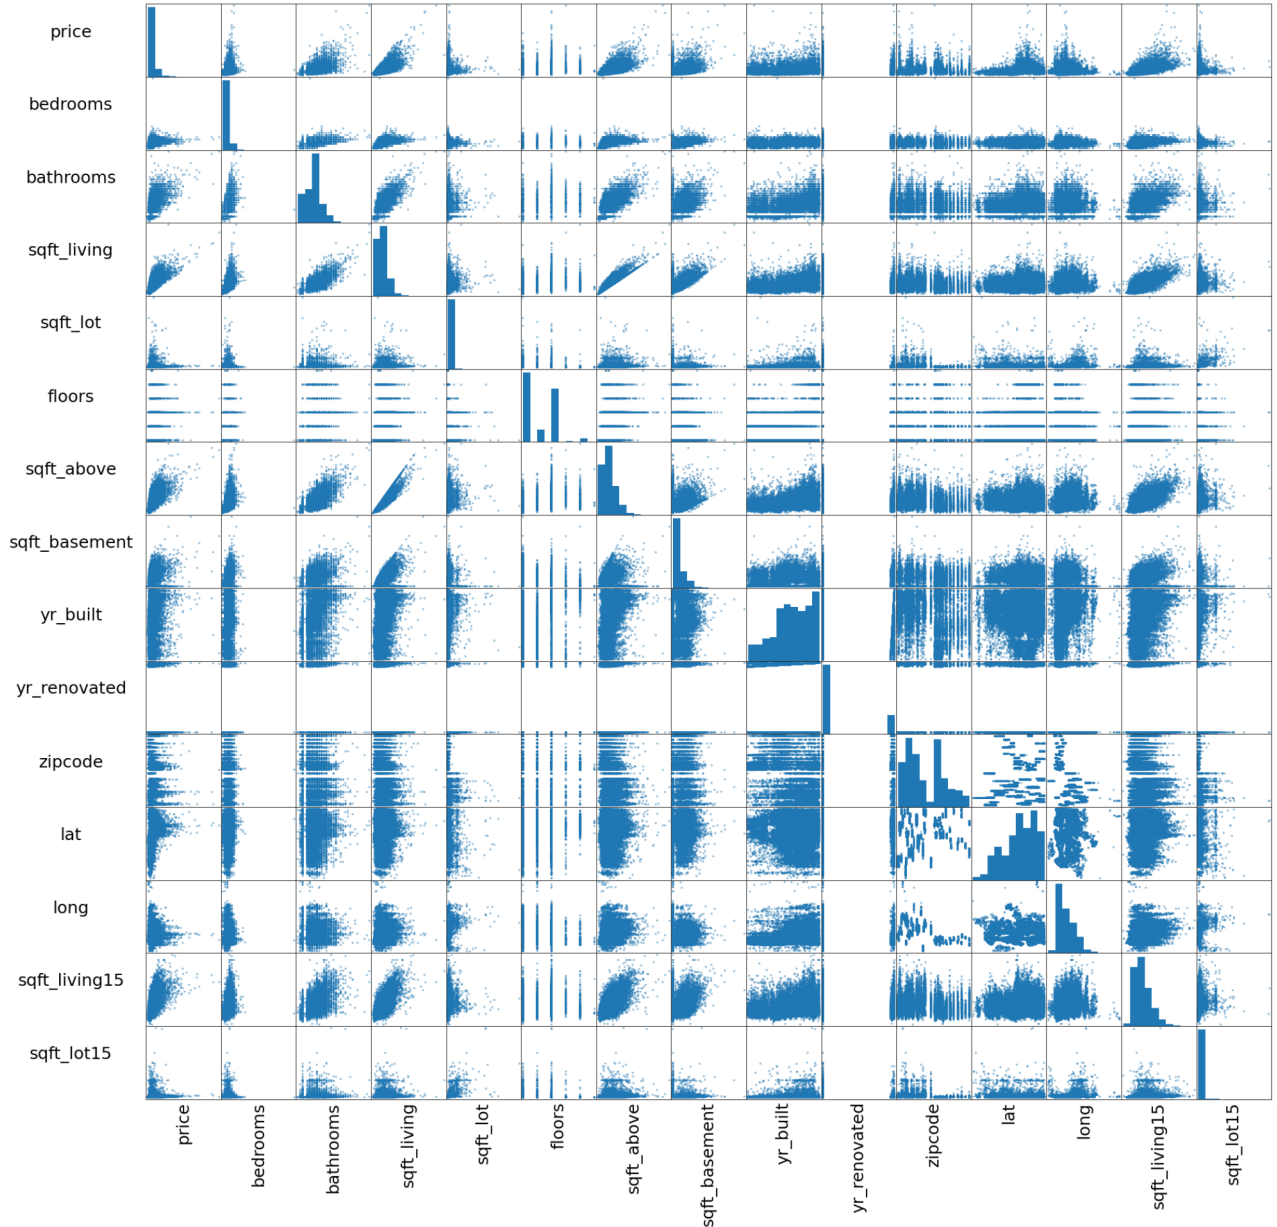
```
In [26]: # create scatter matrix
smat = pd.plotting.scatter_matrix(df, figsize=[30, 30]);

# Rotates the text
[s.xaxis.label.set_rotation(90) for s in smat.reshape(-1)]
[s.yaxis.label.set_rotation(0) for s in smat.reshape(-1)]

#May need to offset label when rotating to prevent overlap of figure
[s.get_yaxis().set_label_coords(-1,0.5) for s in smat.reshape(-1)]


#Hide all ticks
[s.set_xticks(()) for s in smat.reshape(-1)]
[s.set_yticks(()) for s in smat.reshape(-1)]
[plt.setp(item.xaxis.get_label(), "size", 25) for item in smat.ravel()]
[plt.setp(item.yaxis.get_label(), "size", 25) for item in smat.ravel()]

plt.show()
```




Visualize target using seaborn figure level plot called `displot` .

```
In [27]: sns.set_context("paper", rc={"font.size":12,"axes.labelsize":14})
sns.displot(df['price'], stat = 'density', kde = True, height = 15)
plt.ticklabel_format(style='plain', axis= 'both')
plt.show();
```



The displot shows that most of the houses in the KC House data set are priced below 1 million dollars.

```
In [28]: # Rough filter investigating dataframe upper limit
new_df = df[df['price'] <= 1000000]
sns.set_context("paper", rc={"font.size":12,"axes.labelsize":14})
sns.displot(new_df['price'], stat = 'density', kde = True, height = 15)
plt.ticklabel_format(style='plain', axis= 'both')
plt.show();
```




```
In [29]: # Identify the name of the predictor column with strongest correlation
most_correlated = 'sqft_living'
```

From reviewing dataframe df with correlation coefficient values between 0.6 and 1, we can drop the following features from our original dataframe (df) to create a new dataframe called df_subset :

- sqft_above
- sqft_living15
- bathrooms
- sqft_lot
- sqft_lot15

```
In [30]: df_subset = new_df.drop(["sqft_above", "sqft_living15", "bathrooms",
                                , "sqft_lot", "sqft_lot15"], axis = 1)
```

```
In [31]: plt.style.use('fivethirtyeight')
values = np.array(df_subset.sample(1000, random_state = 1).price)/1000000
# Set up plot
fig, ax = plt.subplots(figsize = (15, 10))
sns.regplot(data = df_subset.sample(1000, random_state = 1), x = most_correlated, y = values
            , scatter = True, ci = 95, fit_reg = True, scatter_kws={"color": "blue"}, line_kws={"color": "black"} );
sns.regplot(data = df_subset.sample(1000, random_state = 1), x = most_correlated, y = values, scatter = False
            , ci = 0, fit_reg = True, color = 'red');
ax.set_title("Square Ft of Living Space vs Price", weight = 'bold').set_fontsize('16');
ax.set_xlabel("sqft_living (sqare feet)", fontsize = '14', weight = 'bold');
ax.set_ylabel("Price ($1M dollars)", fontsize = '14', weight = 'bold');
```



6 Build a Simple Linear Regression model

First, I will set the dependent variable (y) to be the `price` . Next I will choose the most highly correlated features from the dataframe to be the baseline independent variable (x). Afterwards, I will:

- Build a linear regression using `statsModels`
- Describe the overall model performance
- Interpret its coefficients.

```
In [32]: # Explore correlation to find a good starting point
df_subset.corr()['price'].sort_values()
```

```
Out[32]: zipcode      -0.02
yr_renovated    0.03
yr_built        0.06
long            0.08
sqft_basement   0.22
floors          0.27
bedrooms        0.29
lat             0.44
sqft_living     0.60
price           1.00
Name: price, dtype: float64
```

```
In [33]: # Set price as the dependent variable
y = df_subset["price"]
```

```
In [34]: # sqft_living had the highest correlation
X_baseline = df_subset[[most_correlated]]
```

6.1 Creating and Fitting Simple Linear Regression

```
In [35]: baseline_model = sm.OLS(y, sm.add_constant(X_baseline))
baseline_results = baseline_model.fit()
```

7 Evaluate and Interpret Baseline Model Results

```
In [36]: print(baseline_results.summary())
```

```

=====
                        OLS Regression Results
=====
Dep. Variable:          price      R-squared:                0.365
Model:                  OLS      Adj. R-squared:             0.365
Method:                 Least Squares    F-statistic:         1.158e+04
Date:                   Mon, 16 Jan 2023    Prob (F-statistic):       0.00
Time:                   21:01:15    Log-Likelihood:        -2.6946e+05
No. Observations:       20139    AIC:                   5.389e+05
Df Residuals:           20137    BIC:                   5.389e+05
Df Model:                1
Covariance Type:        nonrobust
=====
                        coef      std err          t      P>|t|      [0.025      0.975]
-----
const          1.607e+05    3061.929     52.468     0.000     1.55e+05     1.67e+05
sqft_living    157.1496         1.461    107.596     0.000     154.287     160.012
=====
Omnibus:                 704.434    Durbin-Watson:           1.961
Prob(Omnibus):            0.000    Jarque-Bera (JB):         780.694
Skew:                     0.482    Prob(JB):                 2.98e-170
Kurtosis:                 2.973    Cond. No.                  5.82e+03
=====

Notes:
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
[2] The condition number is large, 5.82e+03. This might indicate that there are
strong multicollinearity or other numerical problems.
```

This simple linear regression model is statistically significant overall, and explains **49.3% of the variance in house price**. Both the intercept and the coefficient for `sqft_living` are statistically significant.

The intercept is a small negative number, meaning a home with 0 square feet of living would cost around \$0.

The coefficient for `sqft_living` is about 280, which means that for each additional square foot of living space, I expect the price to increase about **\$280**.

8 Prepare Categorical Features for Multiple Regression Modeling

The categorical features that need to be replaced with dummies are `grade`, `view`, `waterfront`, and `zipcode`. First, let's review how many values each of these features has.

```
In [37]: new_df["grade"].value_counts()
```

```
Out[37]: 7 Average          8951
8 Good              5873
9 Better            2224
6 Low Average       2033
10 Very Good        694
5 Fair              242
11 Excellent         92
4 Low                27
12 Luxury            2
3 Poor               1
Name: grade, dtype: int64
```

```
In [38]: new_df["view"].value_counts()
```

```
Out[38]: NONE          18611
AVERAGE         763
GOOD             319
FAIR             260
EXCELLENT        129
0                 57
Name: view, dtype: int64
```

```
In [39]: new_df["waterfront"].value_counts()
```

```
Out[39]: NO          17862
0           2227
YES          50
Name: waterfront, dtype: int64
```

```
In [40]: new_df["zipcode"].value_counts()
```

```
Out[40]: 98038      586
98103      581
98052      553
98115      549
98042      547
...
98109       82
98102       81
98024       72
98148       57
98039        6
Name: zipcode, Length: 70, dtype: int64
```

Next, lets make a new dataframe that includes these categoricals called `df_updated` .

```
In [41]: df_updated = new_df[['bedrooms', 'sqft_living', 'lat', 'long', "grade", "view", "zipcode", "waterfront"]]
df_updated.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 20139 entries, 7129300520 to 1523300157
Data columns (total 8 columns):
#   Column          Non-Null Count  Dtype
---  -
0   bedrooms        20139 non-null  int64
1   sqft_living      20139 non-null  int64
2   lat              20139 non-null  float64
3   long             20139 non-null  float64
4   grade            20139 non-null  object
5   view             20139 non-null  object
6   zipcode          20139 non-null  int64
7   waterfront       20139 non-null  object
dtypes: float64(2), int64(3), object(3)
memory usage: 1.4+ MB
```

Now let's use **dummy variables** to make this dataframe usable for a linnear regression model.

```
In [42]: # Create a new data frame with view dummy variables
df_with_dummies = pd.get_dummies(data = df_updated, columns = ["view", "zipcode", "grade", "waterfront"], drop_first = True)
```

```
In [43]: df_with_dummies.head()
```

```
Out[43]:
```

	bedrooms	sqft_living	lat	long	view_AVERAGE	view_EXCELLENT	view_FAIR	view_GOOD	view_NONE	zipcode_98002	...	grade_12	Luxury
id													
7129300520	3	1180	47.51	-122.26	0	0	0	0	1	0	...	0	
6414100192	3	2570	47.72	-122.32	0	0	0	0	1	0	...	0	
5631500400	2	770	47.74	-122.23	0	0	0	0	1	0	...	0	
2487200875	4	1960	47.52	-122.39	0	0	0	0	1	0	...	0	
1954400510	3	1680	47.62	-122.05	0	0	0	0	1	0	...	0	

5 rows × 89 columns


```
In [44]: target = new_df["price"]
```

```
In [45]: target
```

```
Out[45]: id
7129300520    221900.00
6414100192    538000.00
5631500400    180000.00
2487200875    604000.00
1954400510    510000.00
...
263000018    360000.00
6600060120    400000.00
1523300141    402101.00
291310100     400000.00
1523300157    325000.00
Name: price, Length: 20139, dtype: float64
```

```
In [46]: iterated_model = sm.OLS(target, sm.add_constant(df_with_dummies))
iterated_results = iterated_model.fit()
```

```
In [47]: print(iterated_results.summary())
```

```

                    OLS Regression Results
=====
Dep. Variable:      price    R-squared:                0.813
Model:              OLS      Adj. R-squared:            0.812
Method:             Least Squares    F-statistic:        978.0
Date:               Mon, 16 Jan 2023    Prob (F-statistic):    0.00
Time:                21:01:15      Log-Likelihood:      -2.5717e+05
No. Observations:    20139      AIC:                  5.145e+05
Df Residuals:        20049      BIC:                  5.152e+05
Df Model:             89
Covariance Type:     nonrobust
=====
                    coef    std err          t      P>|t|      [0.025     0.975]
-----
const             -1.142e+07   3.34e+06   -3.419    0.001   -1.8e+07   -4.87e+06
bedrooms           -428.4762    860.000   -0.498    0.618   -2114.147   1257.195
sqft_living         109.9166     1.387    79.261    0.000    107.198    112.635
lat                1.465e+05    3.47e+04    4.223    0.000    7.85e+04    2.14e+05
long              -3.85e+04    2.47e+04   -1.560    0.119   -8.69e+04    9863.715
view_AVERAGE       1.512e+04    1.17e+04    1.273    0.200    2.21e+03   2.80e+04
view_FAIR           1.512e+04    1.17e+04    1.273    0.200    2.21e+03   2.80e+04
zipcode_98002       1.512e+04    1.17e+04    1.273    0.200    2.21e+03   2.80e+04
zipcode_98003       1.512e+04    1.17e+04    1.273    0.200    2.21e+03   2.80e+04
zipcode_98011       1.512e+04    1.17e+04    1.273    0.200    2.21e+03   2.80e+04
zipcode_98014       1.512e+04    1.17e+04    1.273    0.200    2.21e+03   2.80e+04
zipcode_98077       1.512e+04    1.17e+04    1.273    0.200    2.21e+03   2.80e+04
zipcode_98028       1.512e+04    1.17e+04    1.273    0.200    2.21e+03   2.80e+04
zipcode_98030       1.512e+04    1.17e+04    1.273    0.200    2.21e+03   2.80e+04
zipcode_98031       1.512e+04    1.17e+04    1.273    0.200    2.21e+03   2.80e+04
zipcode_98032       1.512e+04    1.17e+04    1.273    0.200    2.21e+03   2.80e+04
zipcode_98042       1.512e+04    1.17e+04    1.273    0.200    2.21e+03   2.80e+04
zipcode_98055       1.512e+04    1.17e+04    1.273    0.200    2.21e+03   2.80e+04
zipcode_98058       1.512e+04    1.17e+04    1.273    0.200    2.21e+03   2.80e+04
zipcode_98077       1.512e+04    1.17e+04    1.273    0.200    2.21e+03   2.80e+04
zipcode_98188       1.512e+04    1.17e+04    1.273    0.200    2.21e+03   2.80e+04
zipcode_98148       1.512e+04    1.17e+04    1.273    0.200    2.21e+03   2.80e+04
zipcode_98168       1.512e+04    1.17e+04    1.273    0.200    2.21e+03   2.80e+04
zipcode_98092       1.512e+04    1.17e+04    1.273    0.200    2.21e+03   2.80e+04
zipcode_98198       1.512e+04    1.17e+04    1.273    0.200    2.21e+03   2.80e+04
zipcode_98070       1.512e+04    1.17e+04    1.273    0.200    2.21e+03   2.80e+04
grade_3_Poor       1.512e+04    1.17e+04    1.273    0.200    2.21e+03   2.80e+04
waterfront_NO      1.512e+04    1.17e+04    1.273    0.200    2.21e+03   2.80e+04
=====
```

8.1 Remove the values with high p-values

```
In [48]: revised_df = df_with_dummies.drop(["view_AVERAGE", "view_FAIR", "zipcode_98002"
, "zipcode_98003", "zipcode_98011", "zipcode_98014"
, "zipcode_98077", "zipcode_98028", "zipcode_98030"
, "zipcode_98031", "zipcode_98032", "zipcode_98042"
, "zipcode_98055", "zipcode_98058", "zipcode_98077", "zipcode_98188"
, "zipcode_98148", "zipcode_98168", "zipcode_98092", "zipcode_98198"
, "zipcode_98070", "grade_3_Poor", "waterfront_NO"], axis = 1)
```

```
In [73]: revised_model = sm.OLS(target, sm.add_constant(revised_df))
revised_results = revised_model.fit()
```

```
In [74]: ► print(revised_results.summary())
```

OLS Regression Results

```

=====
Dep. Variable:          price      R-squared:                0.810
Model:                  OLS      Adj. R-squared:             0.809
Method:                 Least Squares      F-statistic:          1274.
Date:                   Mon, 16 Jan 2023   Prob (F-statistic):      0.00
Time:                   21:14:03          Log-Likelihood:         -2.5734e+05
No. Observations:      20139            AIC:                   5.148e+05
Df Residuals:          20071            BIC:                   5.153e+05
Df Model:               67
Covariance Type:       nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
const	-1.736e+07	1.55e+06	-11.190	0.000	-2.04e+07	-1.43e+07
bedrooms	-1100.0639	864.564	-1.272	0.203	-2794.680	594.553
sqft_living	110.9095	1.395	79.485	0.000	108.174	113.644
lat	3.45e+05	8638.089	39.943	0.000	3.28e+05	3.62e+05
long	-1.065e+04	1.16e+04	-0.920	0.358	-3.33e+04	1.2e+04
view_EXCELLENT	9.417e+04	9003.015	10.460	0.000	7.65e+04	1.12e+05
view_GOOD	2.571e+04	5505.471	4.669	0.000	1.49e+04	3.65e+04
view_NONE	-6.075e+04	2779.500	-21.858	0.000	-6.62e+04	-5.53e+04
zipcode_98004	4.075e+05	7588.574	53.699	0.000	3.93e+05	4.22e+05
zipcode_98005	2.463e+05	7371.820	33.414	0.000	2.32e+05	2.61e+05
zipcode_98006	1.987e+05	4725.756	42.041	0.000	1.89e+05	2.08e+05
zipcode_98007	1.649e+05	7683.570	21.455	0.000	1.5e+05	1.8e+05
zipcode_98008	1.513e+05	5694.763	26.569	0.000	1.4e+05	1.62e+05
zipcode_98010	9.98e+04	9111.519	10.953	0.000	8.19e+04	1.18e+05
zipcode_98019	-3.58e+04	7246.918	-4.940	0.000	-5e+04	-2.16e+04
zipcode_98022	6.779e+04	6769.352	10.014	0.000	5.45e+04	8.11e+04
zipcode_98023	-1.83e+04	4495.090	-4.071	0.000	-2.71e+04	-9490.531
zipcode_98024	9.049e+04	1.08e+04	8.397	0.000	6.94e+04	1.12e+05
zipcode_98027	1.297e+05	4906.830	26.437	0.000	1.2e+05	1.39e+05
zipcode_98029	1.443e+05	5598.263	25.779	0.000	1.33e+05	1.55e+05
zipcode_98033	1.932e+05	5219.876	37.007	0.000	1.83e+05	2.03e+05
zipcode_98034	4.701e+04	4654.291	10.101	0.000	3.79e+04	5.61e+04
zipcode_98038	2.925e+04	4376.620	6.682	0.000	2.07e+04	3.78e+04
zipcode_98039	5.598e+05	3.51e+04	15.936	0.000	4.91e+05	6.29e+05
zipcode_98040	3.479e+05	7302.947	47.638	0.000	3.34e+05	3.62e+05
zipcode_98045	6.236e+04	7893.578	7.900	0.000	4.69e+04	7.78e+04
zipcode_98052	1.361e+05	4407.164	30.880	0.000	1.27e+05	1.45e+05
zipcode_98053	1.245e+05	5343.599	23.295	0.000	1.14e+05	1.35e+05
zipcode_98056	3.89e+04	4518.904	8.609	0.000	3e+04	4.78e+04
zipcode_98059	5.175e+04	4361.658	11.864	0.000	4.32e+04	6.03e+04
zipcode_98065	6.351e+04	6434.366	9.870	0.000	5.09e+04	7.61e+04
zipcode_98072	3.898e+04	6071.862	6.421	0.000	2.71e+04	5.09e+04
zipcode_98074	1.155e+05	4954.809	23.307	0.000	1.06e+05	1.25e+05
zipcode_98075	1.573e+05	5560.670	28.293	0.000	1.46e+05	1.68e+05
zipcode_98102	2.692e+05	9928.257	27.119	0.000	2.5e+05	2.89e+05
zipcode_98103	1.877e+05	4731.373	39.666	0.000	1.78e+05	1.97e+05
zipcode_98105	2.555e+05	7113.987	35.911	0.000	2.42e+05	2.69e+05
zipcode_98106	2.181e+04	5294.405	4.120	0.000	1.14e+04	3.22e+04
zipcode_98107	1.888e+05	6270.634	30.116	0.000	1.77e+05	2.01e+05
zipcode_98108	2.576e+04	6626.665	3.887	0.000	1.28e+04	3.87e+04
zipcode_98109	2.914e+05	9922.805	29.370	0.000	2.72e+05	3.11e+05
zipcode_98112	2.993e+05	7421.443	40.329	0.000	2.85e+05	3.14e+05
zipcode_98115	1.88e+05	4660.527	40.344	0.000	1.79e+05	1.97e+05
zipcode_98116	1.962e+05	5723.631	34.270	0.000	1.85e+05	2.07e+05
zipcode_98117	1.769e+05	5056.509	34.991	0.000	1.67e+05	1.87e+05
zipcode_98118	7.413e+04	4271.551	17.353	0.000	6.58e+04	8.25e+04
zipcode_98119	2.776e+05	7731.911	35.899	0.000	2.62e+05	2.93e+05
zipcode_98122	2.003e+05	5760.419	34.776	0.000	1.89e+05	2.12e+05
zipcode_98125	5.19e+04	5292.516	9.806	0.000	4.15e+04	6.23e+04
zipcode_98126	1.005e+05	5291.665	18.983	0.000	9.01e+04	1.11e+05
zipcode_98133	3779.5949	5221.256	0.724	0.469	-6454.496	1.4e+04
zipcode_98136	1.635e+05	6074.849	26.913	0.000	1.52e+05	1.75e+05
zipcode_98144	1.342e+05	5347.013	25.095	0.000	1.24e+05	1.45e+05
zipcode_98146	3.219e+04	5630.675	5.717	0.000	2.12e+04	4.32e+04
zipcode_98155	-1.072e+04	5314.075	-2.017	0.044	-2.11e+04	-301.328
zipcode_98166	5.433e+04	5892.126	9.221	0.000	4.28e+04	6.59e+04
zipcode_98177	7.358e+04	6822.137	10.785	0.000	6.02e+04	8.69e+04
zipcode_98178	-2.26e+04	5557.390	-4.067	0.000	-3.35e+04	-1.17e+04
zipcode_98199	2.256e+05	6371.699	35.413	0.000	2.13e+05	2.38e+05
grade_11 Excellent	5.693e+04	9570.733	5.948	0.000	3.82e+04	7.57e+04
grade_12 Luxury	-2.176e+05	6.11e+04	-3.561	0.000	-3.37e+05	-9.78e+04
grade_4 Low	-1.767e+05	1.72e+04	-10.285	0.000	-2.1e+05	-1.43e+05
grade_5 Fair	-1.836e+05	7084.645	-25.918	0.000	-1.98e+05	-1.7e+05
grade_6 Low Average	-1.76e+05	4683.078	-37.575	0.000	-1.85e+05	-1.67e+05
grade_7 Average	-1.543e+05	4057.989	-38.012	0.000	-1.62e+05	-1.46e+05
grade_8 Good	-1.173e+05	3822.404	-30.684	0.000	-1.25e+05	-1.1e+05
grade_9 Better	-4.505e+04	3850.252	-11.701	0.000	-5.26e+04	-3.75e+04
waterfront_YES	1.851e+05	1.4e+04	13.260	0.000	1.58e+05	2.13e+05

```

=====
Omnibus:                  1674.547      Durbin-Watson:              1.985
Prob(Omnibus):            0.000        Jarque-Bera (JB):           5577.670
Skew:                     0.406         Prob(JB):                   0.00

```

Kurtosis: 5.447 Cond. No. 5.38e+06
=====

Notes:


[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
[2] The condition number is large, 5.38e+06. This might indicate that there are strong multicollinearity or other numerical problems.

```
In [51]: ► baseline_results.rsquared_adj, revised_results.rsquared_adj
```

```
Out[51]: (0.365011142951163, 0.8089660021397984)
```

```
In [79]: ► # check for the normality of the residuals  
sm.qqplot(iterated_results.resid, line='s')  
# also check that the mean of the residuals is approx. 0.  
mean_residuals = sum(iterated_results.resid) / len(iterated_results.resid)  
print("The mean of the residuals is {:.4}".format(mean_residuals))
```

The mean of the residuals is 1.99e-07



```
In [52]: M revised_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 20139 entries, 7129300520 to 1523300157
Data columns (total 67 columns):
#   Column                Non-Null Count  Dtype
---  -
0   bedrooms              20139 non-null  int64
1   sqft_living           20139 non-null  int64
2   lat                   20139 non-null  float64
3   long                  20139 non-null  float64
4   view_EXCELLENT        20139 non-null  uint8
5   view_GOOD             20139 non-null  uint8
6   view_NONE             20139 non-null  uint8
7   zipcode_98004         20139 non-null  uint8
8   zipcode_98005         20139 non-null  uint8
9   zipcode_98006         20139 non-null  uint8
10  zipcode_98007         20139 non-null  uint8
11  zipcode_98008         20139 non-null  uint8
12  zipcode_98010         20139 non-null  uint8
13  zipcode_98019         20139 non-null  uint8
14  zipcode_98022         20139 non-null  uint8
15  zipcode_98023         20139 non-null  uint8
16  zipcode_98024         20139 non-null  uint8
17  zipcode_98027         20139 non-null  uint8
18  zipcode_98029         20139 non-null  uint8
19  zipcode_98033         20139 non-null  uint8
20  zipcode_98034         20139 non-null  uint8
21  zipcode_98038         20139 non-null  uint8
22  zipcode_98039         20139 non-null  uint8
23  zipcode_98040         20139 non-null  uint8
24  zipcode_98045         20139 non-null  uint8
25  zipcode_98052         20139 non-null  uint8
26  zipcode_98053         20139 non-null  uint8
27  zipcode_98056         20139 non-null  uint8
28  zipcode_98059         20139 non-null  uint8
29  zipcode_98065         20139 non-null  uint8
30  zipcode_98072         20139 non-null  uint8
31  zipcode_98074         20139 non-null  uint8
32  zipcode_98075         20139 non-null  uint8
33  zipcode_98102         20139 non-null  uint8
34  zipcode_98103         20139 non-null  uint8
35  zipcode_98105         20139 non-null  uint8
36  zipcode_98106         20139 non-null  uint8
37  zipcode_98107         20139 non-null  uint8
38  zipcode_98108         20139 non-null  uint8
39  zipcode_98109         20139 non-null  uint8
40  zipcode_98112         20139 non-null  uint8
41  zipcode_98115         20139 non-null  uint8
42  zipcode_98116         20139 non-null  uint8
43  zipcode_98117         20139 non-null  uint8
44  zipcode_98118         20139 non-null  uint8
45  zipcode_98119         20139 non-null  uint8
46  zipcode_98122         20139 non-null  uint8
47  zipcode_98125         20139 non-null  uint8
48  zipcode_98126         20139 non-null  uint8
49  zipcode_98133         20139 non-null  uint8
50  zipcode_98136         20139 non-null  uint8
51  zipcode_98144         20139 non-null  uint8
52  zipcode_98146         20139 non-null  uint8
53  zipcode_98155         20139 non-null  uint8
54  zipcode_98166         20139 non-null  uint8
55  zipcode_98177         20139 non-null  uint8
56  zipcode_98178         20139 non-null  uint8
57  zipcode_98199         20139 non-null  uint8
58  grade_11 Excellent    20139 non-null  uint8
59  grade_12 Luxury       20139 non-null  uint8
60  grade_4 Low           20139 non-null  uint8
61  grade_5 Fair          20139 non-null  uint8
62  grade_6 Low Average   20139 non-null  uint8
63  grade_7 Average       20139 non-null  uint8
64  grade_8 Good          20139 non-null  uint8
65  grade_9 Better        20139 non-null  uint8
66  waterfront_YES        20139 non-null  uint8
dtypes: float64(2), int64(2), uint8(63)
memory usage: 2.0 MB
```

```
In [53]: final_df = pd.concat([target, revised_df], axis=1)
final_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 20139 entries, 7129300520 to 1523300157
Data columns (total 68 columns):
#   Column              Non-Null Count  Dtype
---  -
0   price               20139 non-null  float64
1   bedrooms            20139 non-null  int64
2   sqft_living         20139 non-null  int64
3   lat                 20139 non-null  float64
4   long                20139 non-null  float64
5   view_EXCELLENT      20139 non-null  uint8
6   view_GOOD           20139 non-null  uint8
7   view_NONE           20139 non-null  uint8
8   zipcode_98004       20139 non-null  uint8
9   zipcode_98005       20139 non-null  uint8
10  zipcode_98006       20139 non-null  uint8
11  zipcode_98007       20139 non-null  uint8
12  zipcode_98008       20139 non-null  uint8
13  zipcode_98010       20139 non-null  uint8
14  zipcode_98019       20139 non-null  uint8
15  zipcode_98022       20139 non-null  uint8
16  zipcode_98023       20139 non-null  uint8
17  zipcode_98024       20139 non-null  uint8
18  zipcode_98027       20139 non-null  uint8
19  zipcode_98029       20139 non-null  uint8
20  zipcode_98033       20139 non-null  uint8
21  zipcode_98034       20139 non-null  uint8
22  zipcode_98038       20139 non-null  uint8
23  zipcode_98039       20139 non-null  uint8
24  zipcode_98040       20139 non-null  uint8
25  zipcode_98045       20139 non-null  uint8
26  zipcode_98052       20139 non-null  uint8
27  zipcode_98053       20139 non-null  uint8
28  zipcode_98056       20139 non-null  uint8
29  zipcode_98059       20139 non-null  uint8
30  zipcode_98065       20139 non-null  uint8
31  zipcode_98072       20139 non-null  uint8
32  zipcode_98074       20139 non-null  uint8
33  zipcode_98075       20139 non-null  uint8
34  zipcode_98102       20139 non-null  uint8
35  zipcode_98103       20139 non-null  uint8
36  zipcode_98105       20139 non-null  uint8
37  zipcode_98106       20139 non-null  uint8
38  zipcode_98107       20139 non-null  uint8
39  zipcode_98108       20139 non-null  uint8
40  zipcode_98109       20139 non-null  uint8
41  zipcode_98112       20139 non-null  uint8
42  zipcode_98115       20139 non-null  uint8
43  zipcode_98116       20139 non-null  uint8
44  zipcode_98117       20139 non-null  uint8
45  zipcode_98118       20139 non-null  uint8
46  zipcode_98119       20139 non-null  uint8
47  zipcode_98122       20139 non-null  uint8
48  zipcode_98125       20139 non-null  uint8
49  zipcode_98126       20139 non-null  uint8
50  zipcode_98133       20139 non-null  uint8
51  zipcode_98136       20139 non-null  uint8
52  zipcode_98144       20139 non-null  uint8
53  zipcode_98146       20139 non-null  uint8
54  zipcode_98155       20139 non-null  uint8
55  zipcode_98166       20139 non-null  uint8
56  zipcode_98177       20139 non-null  uint8
57  zipcode_98178       20139 non-null  uint8
58  zipcode_98199       20139 non-null  uint8
59  grade_11 Excellent  20139 non-null  uint8
60  grade_12 Luxury     20139 non-null  uint8
61  grade_4 Low         20139 non-null  uint8
62  grade_5 Fair        20139 non-null  uint8
63  grade_6 Low Average 20139 non-null  uint8
64  grade_7 Average     20139 non-null  uint8
65  grade_8 Good        20139 non-null  uint8
66  grade_9 Better      20139 non-null  uint8
67  waterfront_YES      20139 non-null  uint8
dtypes: float64(3), int64(2), uint8(63)
memory usage: 2.1 MB
```

```
In [54]: ► bedrooms_counts = final_df['bedrooms'].value_counts()  
print(bedrooms_counts)
```

```
3    9519  
4    6147  
2    2727  
5    1286  
6     218  
1     195  
7       30  
8        9  
9        4  
10       2  
11       1  
33       1  
Name: bedrooms, dtype: int64
```

```
In [55]: # Remove Bedrooms that are 9 or larger from final dataframe
final_df = final_df[final_df['bedrooms'] < 9]
final_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 20131 entries, 7129300520 to 1523300157
Data columns (total 68 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   price                                20131 non-null  float64
1   bedrooms                            20131 non-null  int64
2   sqft_living                         20131 non-null  int64
3   lat                                 20131 non-null  float64
4   long                                20131 non-null  float64
5   view_EXCELLENT                      20131 non-null  uint8
6   view_GOOD                           20131 non-null  uint8
7   view_NONE                           20131 non-null  uint8
8   zipcode_98004                       20131 non-null  uint8
9   zipcode_98005                       20131 non-null  uint8
10  zipcode_98006                       20131 non-null  uint8
11  zipcode_98007                       20131 non-null  uint8
12  zipcode_98008                       20131 non-null  uint8
13  zipcode_98010                       20131 non-null  uint8
14  zipcode_98019                       20131 non-null  uint8
15  zipcode_98022                       20131 non-null  uint8
16  zipcode_98023                       20131 non-null  uint8
17  zipcode_98024                       20131 non-null  uint8
18  zipcode_98027                       20131 non-null  uint8
19  zipcode_98029                       20131 non-null  uint8
20  zipcode_98033                       20131 non-null  uint8
21  zipcode_98034                       20131 non-null  uint8
22  zipcode_98038                       20131 non-null  uint8
23  zipcode_98039                       20131 non-null  uint8
24  zipcode_98040                       20131 non-null  uint8
25  zipcode_98045                       20131 non-null  uint8
26  zipcode_98052                       20131 non-null  uint8
27  zipcode_98053                       20131 non-null  uint8
28  zipcode_98056                       20131 non-null  uint8
29  zipcode_98059                       20131 non-null  uint8
30  zipcode_98065                       20131 non-null  uint8
31  zipcode_98072                       20131 non-null  uint8
32  zipcode_98074                       20131 non-null  uint8
33  zipcode_98075                       20131 non-null  uint8
34  zipcode_98102                       20131 non-null  uint8
35  zipcode_98103                       20131 non-null  uint8
36  zipcode_98105                       20131 non-null  uint8
37  zipcode_98106                       20131 non-null  uint8
38  zipcode_98107                       20131 non-null  uint8
39  zipcode_98108                       20131 non-null  uint8
40  zipcode_98109                       20131 non-null  uint8
41  zipcode_98112                       20131 non-null  uint8
42  zipcode_98115                       20131 non-null  uint8
43  zipcode_98116                       20131 non-null  uint8
44  zipcode_98117                       20131 non-null  uint8
45  zipcode_98118                       20131 non-null  uint8
46  zipcode_98119                       20131 non-null  uint8
47  zipcode_98122                       20131 non-null  uint8
48  zipcode_98125                       20131 non-null  uint8
49  zipcode_98126                       20131 non-null  uint8
50  zipcode_98133                       20131 non-null  uint8
51  zipcode_98136                       20131 non-null  uint8
52  zipcode_98144                       20131 non-null  uint8
53  zipcode_98146                       20131 non-null  uint8
54  zipcode_98155                       20131 non-null  uint8
55  zipcode_98166                       20131 non-null  uint8
56  zipcode_98177                       20131 non-null  uint8
57  zipcode_98178                       20131 non-null  uint8
58  zipcode_98199                       20131 non-null  uint8
59  grade_11 Excellent                  20131 non-null  uint8
60  grade_12 Luxury                     20131 non-null  uint8
61  grade_4 Low                         20131 non-null  uint8
62  grade_5 Fair                        20131 non-null  uint8
63  grade_6 Low Average                 20131 non-null  uint8
64  grade_7 Average                     20131 non-null  uint8
65  grade_8 Good                        20131 non-null  uint8
66  grade_9 Better                      20131 non-null  uint8
67  waterfront_YES                     20131 non-null  uint8
dtypes: float64(3), int64(2), uint8(63)
memory usage: 2.1 MB
```



```
In [56]: ► bedrooms_counts = final_df['bedrooms'].value_counts()
print(bedrooms_counts)
```

```
3    9519
4    6147
2    2727
5    1286
6     218
1     195
7      30
8       9
Name: bedrooms, dtype: int64
```

```
In [57]: ► final_df.describe()
```

Out[57]:

	price	bedrooms	sqft_living	lat	long	view_EXCELLENT	view_GOOD	view_NONE	zipcode_98004	zipcode_98005	...	zipcode_98119
count	20131.00	20131.00	20131.00	20131.00	20131.00	20131.00	20131.00	20131.00	20131.00	20131.00	...	20131.00
mean	467900.80	3.32	1955.09	47.56	-122.21	0.01	0.02	0.92	0.01	0.01	...	0.01
std	196484.62	0.88	755.07	0.14	0.14	0.08	0.12	0.26	0.08	0.08	...	0.14
min	78000.00	1.00	370.00	47.16	-122.52	0.00	0.00	0.00	0.00	0.00	...	0.00
25%	314500.00	3.00	1390.00	47.46	-122.33	0.00	0.00	1.00	0.00	0.00	...	0.00
50%	434900.00	3.00	1840.00	47.57	-122.23	0.00	0.00	1.00	0.00	0.00	...	0.00
75%	594000.00	4.00	2410.00	47.68	-122.12	0.00	0.00	1.00	0.00	0.00	...	0.00
max	1000000.00	8.00	7480.00	47.78	-121.31	1.00	1.00	1.00	1.00	1.00	...	1.00

8 rows × 68 columns

8.2 Exploratory Data Analysis Q1

8.3 Which neighborhoods have the highest average home price?

The first thought I had when trying to answer this question was that I should leverage the location of homes in the dataset and use a sequential color palette. The first way that I tried to visualize which neighborhoods have the highest home price was to plot the longitude and latitude against one another so that I could get a quick geographical visual on where the more expensive homes were located.

```
In [58]: ► import geopandas as gpd


def sale_price_map(df):
    """
    Creates a map of sale prices for homes in king county dataset using geopandas

    Inputs:
    df: pandas dataframe with columns 'bedrooms', 'lat', 'long', and 'price'

    Output:
    map of sale prices for homes in king county
    """
    # Set up plot
    plt.figure(figsize = (20,15))
    # Create a GeoDataFrame from the pandas dataframe
    gdf = gpd.GeoDataFrame(df, geometry=gpd.points_from_xy(df.long, df.lat))
    # Set the CRS (coordinate reference system) of the GeoDataFrame
    gdf.crs = {'init': 'epsg:4326'}
    # Create a map of sale prices for homes in king county
    gdf.plot(column='price', cmap='Blues', legend=True)
    plt.title("Sale Prices of Homes in King County")
    plt.xlabel("Longitude")
    plt.ylabel("Latitude")
    plt.show()
```

```
In [59]: sale_price_map(df_subset)
```

<Figure size 1440x1080 with 0 Axes>



This visualization gave a general idea that the neighborhoods in the North Eastern section of King County were the most pricey and the neighborhoods to the South and Western sections of King County were more affordable. I wanted to make connections at the neighborhood level based on client need, so I used `plotly.express` as `px` and created a figure using a `scatter_mapbox` where I could zoom and hover over individual homes in different neighborhoods.

```
In [60]: import plotly.express as px
import pandas as pd

# Load the data
data = df_subset

# create the map
fig = px.scatter_mapbox(data_frame=data,
                        lat='lat', # column in data that contains Latitude
                        lon='lon', # column in data that contains Longitude
                        color='price', # column in data that contains housing prices
                        title='Housing Prices in King County by Neighborhood Location',
                        hover_name='zipcode', # column in data that contains address of the property
                        size='bedrooms', # column in data that contains square footage of the property
                        zoom=9,
                        color_continuous_scale="Blues",
                        mapbox_style='carto-positron')

fig.show()
```

Housing Prices in King County by Neighborhood Location

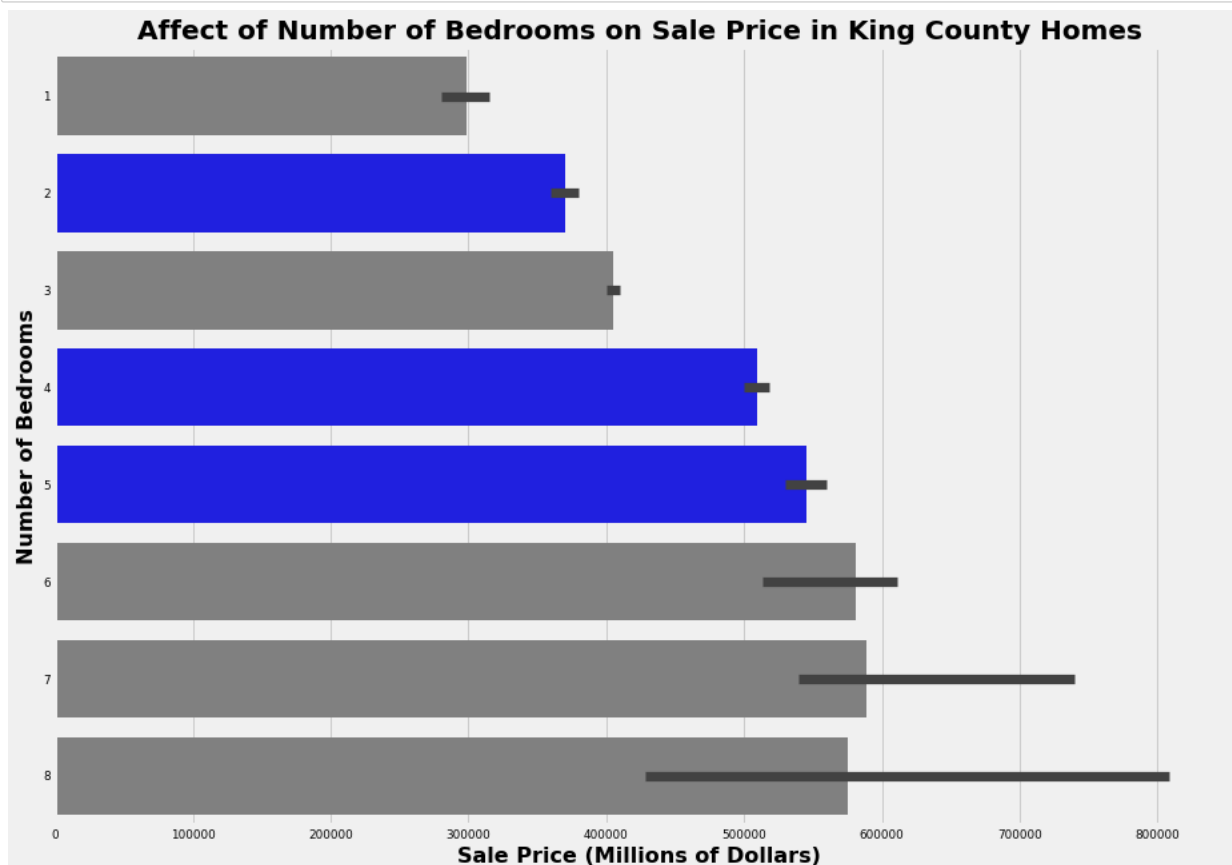


8.4 Exploratory Data Analysis Q2

8.5 How does the number of bedrooms affect the sale price of a home?

```
In [61]: def plot_bedrooms_price(df):  
    """  
    Plot bar graph of number of bedrooms vs price  
  
    Inputs:  
    df: pandas dataframe  
    lim: integer  
    xlabel: independent variable label  
    ylabel: dependent variable label  
  
    Output:  
    bar plot  
    """  
    # Create plot variables  
    prices = np.array(df.price)  
    bedrooms = np.array(df.bedrooms)  
    median_price = np.median(prices)  
  
    # Create plot labels  
    x_label = "Sale Price (Millions of Dollars)"  
    y_label = "Number of Bedrooms"  
    title = "Affect of Number of Bedrooms on Sale Price in King County Homes"  
  
    # Create a palette that highlights the top 5 sale price values as a function of the number of bedrooms"  
    clr = ['grey' if (x < median_price) else 'blue' for x in prices]  
    # Set up plot figure size  
    plt.figure(figsize = (14,10))  
  
    # Plot bar using input of bedrooms, price, index and palette  
    ax = sns.barplot(x = prices, y = bedrooms, palette = clr, estimator = np.median, orient = 'h')  
    # Setup titles and axes labels  
    ax.set_title(title, weight='bold').set_fontsize('20')  
    ax.set_ylabel(y_label, fontsize='16', weight='bold')  
    ax.set_xlabel(x_label, fontsize='16', weight='bold')
```

```
In [62]: plot_bedrooms_price(final_df)
```



```
In [63]: final_df.describe()
```

Out[63]:

	price	bedrooms	sqft_living	lat	long	view_EXCELLENT	view_GOOD	view_NONE	zipcode_98004	zipcode_98005	...	zipcode_9819
count	20131.00	20131.00	20131.00	20131.00	20131.00	20131.00	20131.00	20131.00	20131.00	20131.00	...	20131.00
mean	467900.80	3.32	1955.09	47.56	-122.21	0.01	0.02	0.92	0.01	0.01	...	0.01
std	196484.62	0.88	755.07	0.14	0.14	0.08	0.12	0.26	0.08	0.08	...	0.14
min	78000.00	1.00	370.00	47.16	-122.52	0.00	0.00	0.00	0.00	0.00	...	0.00
25%	314500.00	3.00	1390.00	47.46	-122.33	0.00	0.00	1.00	0.00	0.00	...	0.00
50%	434900.00	3.00	1840.00	47.57	-122.23	0.00	0.00	1.00	0.00	0.00	...	0.00
75%	594000.00	4.00	2410.00	47.68	-122.12	0.00	0.00	1.00	0.00	0.00	...	0.00
max	1000000.00	8.00	7480.00	47.78	-121.31	1.00	1.00	1.00	1.00	1.00	...	1.00

8 rows × 68 columns

8.6 Exploratory Data Analysis Q3

8.7 How does proximity to a highly rated school affect the sale price of a home?

```
In [81]: # Read Tables from the GreatSchools Website
import pandas as pd
from selenium import webdriver

url = "https://www.greatschools.org/search/search.page?lat=47.5480339&locationLabel=King%20County%2C%20WA%2C%20USA&locationTy

driver = webdriver.Chrome('C:\chromedriver_win32\chromedriver.exe')
driver.get(url)

html = driver.page_source

table = pd.read_html(html)
king_county_schools_df = table[0]

driver.close()
```

<ipython-input-81-fb44b494d6b3>:7: DeprecationWarning:

executable_path has been deprecated, please pass in a Service object

```
In [83]: king_county_schools_df.head(20)
```

Out[83]:

		School	Type	Grades	Total students enrolled	Students per teacher	Reviews	District
0	Assigned school7/10	Above averageGrand Ridge El...	Public district	K-5	747	19:1	8 Reviews	Issaquah School District
1	Assigned school7/10	Above averagePacific Cascad...	Public district	6-8	1000	23:1	6 Reviews	Issaquah School District
2	Assigned school8/10	Above averageIssaquah High ...	Public district	9-12	2417	26:1	22 Reviews	Issaquah School District
3	10/10	Above averageSkyline High School2 awardsA...	Public district	9-12	2169	25:1	21 Reviews	Issaquah School District
4	9/10	Above averageSunny Hills Elementary School...	Public district	K-5	732	18:1	11 Reviews	Issaquah School District
5	9/10	Above averageDiscovery Elementary School23...	Public district	PK-5	686	18:1	6 Reviews	Issaquah School District
6	8/10	Above averageFall City Elementary School33...	Public district	K-5	544	19:1	8 Reviews	Snoqualmie Valley School District
7	8/10	Above averageCascade Ridge Elementary Scho...	Public district	K-5	487	17:1	6 Reviews	Issaquah School District
8	8/10	Above averageCreekside Elementary School20...	Public district	K-5	723	18:1	14 Reviews	Issaquah School District
9	8/10	Above averageIssaquah High School2 awardsA...	Public district	9-12	2417	26:1	22 Reviews	Issaquah School District
10	8/10	Above averagePine Lake Middle School3095 l...	Public district	6-8	943	24:1	13 Reviews	Issaquah School District
11	8/10	Above averageEndeavour Elementary School26...	Public district	K-5	598	19:1	9 Reviews	Issaquah School District
12	7/10	Above averageGrand Ridge Elementary School...	Public district	K-5	747	19:1	8 Reviews	Issaquah School District
13	7/10	Above averagePacific Cascade Middle School...	Public district	6-8	1000	23:1	6 Reviews	Issaquah School District
14	7/10	Above averageIssaquah Middle School600 2nd...	Public district	6-8	988	22:1	13 Reviews	Issaquah School District
15	7/10	Above averageBeaver Lake Middle School2502...	Public district	6-8	852	22:1	8 Reviews	Issaquah School District
16	6/10	AverageIssaquah Valley Elementary School55...	Public district	K-5	620	16:1	13 Reviews	Issaquah School District
17	6/10	AverageClark Elementary School335 1st Aven...	Public district	K-5	739	16:1	14 Reviews	Issaquah School District
18	6/10	AverageChallenger Elementary School25200 S...	Public district	K-5	558	17:1	10 Reviews	Issaquah School District
19	5/10	AverageChief Kanim Middle School32627 Redm...	Public district	6-8	795	21:1	12 Reviews	Snoqualmie Valley School District

```
In [64]: from haversine import haversine

def distance_to_school(lat, long, school_lat, school_long):
    # Create a tuple for the home coordinates
    home_coords = (home_lat, home_long)

    # Create a tuple for the school coordinates
    school_coords = (school_lat, school_long)

    # Use the haversine function to calculate the distance between the two points
    distance = haversine(home_coords, school_coords)

    # Return the distance in miles
    return distance
```

9 Insights

10 Recommendations